


Throughput is flat in resolution — the kernel is embarrassingly parallel



Mandelbrot, max_iter=1000. A per-pixel kernel with no shared state should hold constant Mpix/s as the frame grows; it does. The GPU curves rise slightly because their fixed per-frame costs — dispatch and the readback — amortise over more pixels.